

Code Description of Team Polar Bear

Description of the Dispatching Agent

We use a TD(0) on-policy algorithm to dispatch the orders. The details are as follow.

The values are stored by two hashmaps, namely `grid_values` and `layer_values`, which hash grid ids provided by the competition to numerical values. `Grid_values` record the value of a certain grid while `layer_values` are values of grids around a certain grid. All the values are initialized by 0.

For each driver-order pair, we calculate the expected update by:

$$U = (1 - p) \cdot (r + \gamma^t V' - V), \text{ where}$$

p is the estimated cancel probability: $p = 0.01 \cdot \exp(\frac{\log(20)}{2000} \cdot d)$, d is the `order_driver_distance`;

r is the reward of an order;

γ^t is the discount factor, $t = (\text{order_finish_time} - \text{order_start_time})/600$;

V' is the average of `grid_value` and `layer_values` corresponding to the grid where the destination of the order is located;

V is the average of `grid_value` and `layer_values` corresponding to the grid where the driver is located. Then we use KM algorithm to make assignments between drivers and orders where the weight between a driver and an order is the expected update U above. Before applying the KM algorithm, we prune the drivers too far away from an order and only leave top-k nearest drivers of an order where k is set to 11.

According to the assignment results, we update `grid_values` and `layer_values` respectively. If a driver is successfully assigned to an order, the values will be updated by adding $\alpha \cdot U$ where α is the learning rate. Otherwise, i.e., the driver remains idle, the values will be updated by adding $\alpha \cdot (\gamma - 1) \cdot V$.

We do not use any pre-trained models. The KM algorithm is implemented by C++ and we call the function with a dynamic link library “hungnp.so”.

Description of the Repositioning Agent

It is the default repositioning agent that leaves the drivers staying where they are. No RL techniques have been applied.